Pivotal Command Center

Version 2.2

User Guide

Rev: A03 - July 17, 2014

© 2014 Pivotal Software, Inc.

Copyright

Copyright © 2014 GoPivotal, Inc. All Rights reserved.

GoPivotal, Inc. believes the information in this publication is accurate as of its publication date. The information is subject to change without notice. THE INFORMATION IN THIS PUBLICATION IS PROVIDED "AS IS." GOPIVOTAL, INC. ("Pivotal") MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND WITH RESPECT TO THE INFORMATION IN THIS PUBLICATION, AND SPECIFICALLY DISCLAIMS IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE.

Use, copying, and distribution of any Pivotal software described in this publication requires an applicable software license.

All trademarks used herein are the property of Pivotal or their respective owners.

Use of Open Source

This product may be distributed with open source code, licensed to you in accordance with the applicable open source license. If you would like a copy of any such source code, Pivotal will provide a copy of the source code that is required to be made available in accordance with the applicable open source license. Pivotal may charge reasonable shipping and handling charges for such distribution.

About Pivotal, Inc.

Greenplum transitioned to a new corporate identity (Pivotal, Inc.) in 2013. As a result of this transition, there will be some legacy instances of our former corporate identity (Greenplum) appearing in our products and documentation. If you have any questions or concerns, please do not hesitate to contact us through our web site: http://gopivotal.com/about-pivotal/support.

Contents

Chapter 1. PCC Overview	5
About Pivotal Command Center	
Pivotal Command Center UI and CLI	
Performance Monitor (nmon)	
PostgreSQL Database	
Architectural Overview	
, il di illicolara. Ovorvion	
Chapter 2. PCC Installation Checklist	9
Prerequisites	10
Installation Steps	
Chapter 3. Installing PCC	13
Supported Platforms and Browsers	
Platforms	
Browsers	
PCC Prerequisites	
Package Accessibility	
Installation Instructions	20
Installing Pivotal Command Center	20
Import Packages	22
Import JDK	22
Launching Pivotal Command Center	24
Starting, Stopping, and Restarting Command Center Services	24
Uninstalling Pivotal Command Center	25
Upgrading Pivotal Command Center	26
Chapter 4. Using PCC	27
Overview	28
Status indicators	
Logging In	
Login Screen	29
Settings	
User Management	
Users	
Profiles	32
Passwords	33
Cluster Status Page	34

Configuring and Deploying a Cluster	35
Adding a Cluster	35
Starting the Cluster	41
Initializing and Configuring HAWQ	41
Starting Stopping and Uninstalling the Cluster	43
Starting	43
Stopping	43
Uninstalling	43
Dashboard	44
Cluster Analysis	47
MapReduce Job Monitor	49
Job Details	50
YARN App Monitor	53
HAWQ Query Monitor	54
Topology	55
Topology Actions	55
Logs	57
Filtering Logs	57
Viewing Logs	57
Chapter 5. Creating a YUM EPEL Repository	59
Chapter 6. Command Line Reference	60
Backup and Restore	61
Backup	61
Poetoro	61

Chapter 1 PCC Overview

This section provides a brief overview of Pivotal Command Center.

Topics:

- About Pivotal Command Center
 - Pivotal Command Center UI and CLI
 - Performance Monitor (nmon)
 - PostgreSQL Database
- Architectural Overview

About Pivotal Command Center

Pivotal Command Center (PCC) allows an administrative user to configure, deploy, monitor, and manage one or more Pivotal HD clusters. The Command Center has both a graphical user interface and command-line tools to deploy and configure, monitor, and administer Pivotal HD clusters.

- For UI operations, see Using PCC.
- For command-line tools, see the Pivotal HD Enterprise Installation and Administrator Guide.



This release of Command Center allows administering and monitoring of only Pivotal HD Enterprise 2.x clusters.

PCC provides complete life cycle management for Pivotal HD Clusters by performing the following two main groups of functions:

- Cluster configuration and deployment
- Cluster monitoring and management

These functions are served through a set of RESTful web services that run as a web application on Jetty server on the Command Center admin host. This is called gphdmgr-webservices. This web application stores its metadata and cluster configuration for Pivotal HD cluster nodes and services in the Pivotal Command Center PostgreSQL database. It makes use of a Puppet Server to perform most of its HD cluster installation and configuration. It also has a polling service that retrieves Hadoop metrics from the cluster and stores them in the Command Center PostgreSQL Database at periodic intervals.

Pivotal Command Center UI and CLI

The PCC UI provides the user with a single web-based graphical user interface to configure and deploy, monitor and manage one or more Pivotal HD cluster. This web application is hosted on a Ruby-on-Rails application which presents the status and metrics of the clusters. The system metrics data is gathered by the Performance Monitor (nmon) component. The Command Center UI invokes the APIs to retrieve all Hadoop-specific cluster metrics and status information. This includes the Hadoop metrics that was previously retrieved by the polling service.

PCC provides a command-line interface (CLI) for more advanced users to perform installation, configuration and uninstalls. This tool invokes the APIs to install and configure the various Pivotal HD services. The CLI also provides a way to perform other administrative actions such as starting and stopping clusters. For how to use this CLI, please refer to the Pivotal HD Enterprise Installation and Administrator Guide.

Performance Monitor (nmon)

Pivotal Command Center comes with a Performance Monitor called nmon (for node monitor). This makes use of a highly scalable message passing architecture to gather performance metrics from each node that Command Center monitors. This consists of a master daemon that runs on the Command Center admin host and an daemon that runs on all the cluster nodes that report system metric information to the master. This includes metrics such as CPU, memory, disk I/O and network usage information.

The master on the admin host dumps the system metrics it receives from the agents on the cluster nodes into a PostgreSQL DB. This is then queried by the Command Center UI application to display its cluster analysis graphs.

The agents hosts are deployed throughout the cluster during Pivotal HD cluster deployment itself (see for Using PCC details).

The agents are deployed as services on each host, including on the Pivotal Command Center admin host.

To stop or start the service run the following as root:

```
# service nmon stop
# service nmon start
```

PostgreSQL Database

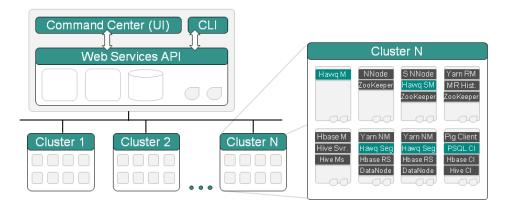
Pivotal Command Center makes use of a PostgreSQL Database to store the following:

- Cluster configurations
- Hadoop cluster metrics
- System metrics of the cluster
- Pivotal Command Center Metadata

Architectural Overview

For more details about Pivotal HD Enterprise, refer to the Pivotal HD Installation and Administrator Guide.

Pivotal Command Center - Architecture



Chapter 2 PCC Installation Checklist

This section is intended to assist you in planning and performing your Pivitol Command Center installation. It includes an installation prerequisite checklist and a brief summary of the installation steps. This section is intended for more advanced users; first time users should following the more detailed instructions we provide in Installing PCC.

Topics:

- Prerequisites
- Installation Steps

Prerequisites

In addition to the specific environment prerequisites, we recommend you have a working knowledge of the following:

- Yum: Yum enables you to install or update software from the command line. See http://yum.baseurl.org/.
- RPM (Redhat Package Manager). See information on RPM at *Managing RPM-Based Systems with Kickstart and Yum.* See http://shop.oreilly.com/product/9780596513825.do?sortby=publicationDate
- NTP. See information on NTP at: http://www.ntp.org
- SSH (Secure Shell Protocol). See information on SSH at http://www.linuxproblem.org/art_9.html

Step	Task	Description	Completed
1	DNS Lookup	Verify that hosts can reach each other using hostnames and IP addresses.	
		<pre>\$ ping -c myhost.mycompany.com // The return code should be 0\$ ping -c 3192.168.1.2 // The return code should be 0</pre>	
2	Check JDK	Ensure you're running Oracle Java JDK Version 1.7 on the Admin node.	
		Java version 1.7 is required; 1.7u15 is recommended	
		As root and gpadmin, run:	
		\$ /usr/java/default/bin/java -version	
		If not, download and install the appropriate version from Oracle.	
3	Package	Verify that all hosts have yum access to an EPEL yum repository.	
	Accessibility	As root, run:	
		<pre>\$ yum list < LIST_OF_PACKAGES ></pre>	
		See Package Accessibility for more details and a list of packages.	
		Note that this is not required if the required rpms are accessible locally.	
4	Verify iptables	As root, run:	
	is turned off	<pre>\$ chkconfig iptables off</pre>	
		\$ service iptables stop	
		<pre>\$ service iptables status</pre>	
		iptables: Firewall is not running.	
5	Disable	As root, run:	
	SELinux	<pre>\$ echo 0 > /selinux/enforce</pre>	

Installation Steps

The table below briefly describes the steps you need to take to upgrade a cluster; more detailed instructions are provided in Installing PCC.

Note that shaded rows depict operations that you perform via a Browser/PCC UI.

Step	Task	Description	Completed
1	Install Pivotal Command Center	As root, run: 1. Create a directory (phd) for your PCC installation: \$ mkdir phd 2. Copy tar file to your specified directory on the admin node, for example: \$ scp ./PCC-2.2.x.version.build.os.x86_64.tar.gz host:/root/phd/ 3. Login as root and untar to that directory: \$ cd /root/phd \$ tarno-same-owner -zxvf PCC-2.2.x.version.build.os.x86_64.tar.gz 4. Run the installation script from the directory where it was extracted: \$./install 5. As the rest of the installation is done as the gpadmin user, change to that user: \$ su - gpadmin 6. If necessary, enable Secure Connections (see Enabling Secure Connections for details)	
2	Import JDK See Import JDK	Download then import a JDK package to the cluster nodes. \$ icm_client import -r <path jdk="" to=""></path>	
3	Copy the PHD Services packages to the Admin Node See Copying PHD Services for more details.	 (as gpadmin) Copy the Pivotal HD services (PHD, ADS (HAWQ)) tarballs from the initial download location to the gpadmin home directory (home/user/gpadmin). Change the owner of the packages to gpadmin then untar the tarballs. For example: If the file is a tar.gz or tgz, use: tar -zxf packagename.tgz If the file is a tar, use: tar -xf packagename.tar 	

Step	Task	Description	Completed
4	Import the PHD services See Import PHD and Import HAWQ for more details.	<pre>(as gpadmin) For each service (PHD, ADS) you are importing, run the following: \$ icm_client import -s < PATH TO EXTRACTED PHD SERVICE TAR BALL ></pre>	
5	Browser: Launch Pivotal Command Center UI See Launching PCC for more details.	Launch a browser and enter the host on which you installed PCC: https://CommandCenterHost:5443 The Command Center login page is launched in your browser. The default username/password is gpadmin/Gpadmin1 (case sensitive).	
6	Browser: Configure and deploy a cluster See Configuring and Deploying a Cluster for more details.	After you have logged in to Pivotal Command Center, the Cluster Status page appears. From here, if you are an administrative user, you are able to launch the Add Cluster Wizard that enables you to configure and deploy a Pivotal HD Cluster.	
7	Browser: Start the cluster See Starting the Cluster for more details.	Return to the PCC UI and start the cluster from the Cluster Status page.	
8	Initialize and Configure HAWQ See Initializing and Configuring HAWQ for more details.	<pre>(as gpadmin) ssh to the HAWQ master, the run the following: \$ source /usr/local/hawq/greenplum_path.sh \$ /etc/init.d/hawq init If you have a HAWQ standby master configured, initialize that: \$ gpinitstandby -s <standby fqdn="" hawq="" master=""></standby></pre>	

Chapter 3 Installing PCC

This section describes how to install Pivotal Command Center and then use PCC to deploy and configure a Pivotal HD cluster.

A PCC Installation Checklist provides a brief overview of the prerequisites and installation instructions; this section provides more detailed information.

This section includes:

Topics:

- Supported Platforms and Browsers
 - Platforms
 - Browsers
- PCC Prerequisites
- Package Accessibility
- Installation Instructions
 - Installing Pivotal Command Center
 - Import Packages
 - Import JDK
 - Launching Pivotal Command Center
 - Starting, Stopping, and Restarting Command Center Services
- Uninstalling Pivotal Command Center
- Upgrading Pivotal Command Center

Supported Platforms and Browsers

Platforms

- RHEL 6.2 64-bit, 6.4 64-bit
- CentOS 6.2 64-bit, 6.4 64-bit

Browsers

(Minimum screen resolution: 1280 x 800)

- Firefox 23
- IE 9, 10 (Chrome Tab not supported)
- Chrome 33.0.1750.146

PCC Prerequisites

Before you begin your installation; we recommend you have working knowledge of the following:

- Yum: Enables you to install or update software from the command line. See http://yum.baseurl.org/.
- RPM (Redhat Package Manager). See information on RPM at Managing RPM-Based Systems with Kickstart and Yum. See http://shop.oreilly.com/product/9780596513825.do?sortby=publicationDate
- NTP. See information on NTP at: http://www.ntp.org
- SSH (Secure Shell Protocol). See information on SSH at http://www.linuxproblem.org/art 9.html

Additionally; the following prerequisites are required:

1. **DNS lookup**. Verify that the admin host (the host on which you will be installing PCC) is be able to reach every host that will be part of your cluster using its hostname and IP address. We also recommend that every cluster node is able to reach every other cluster node using its hostname and IP address:

```
$ ping -c myhost.mycompany.com // The return code should be 0
$ ping -c 3 192.168.1.2 // The return code should be 0
```

2. JAVA JDK. Ensure that you are running Oracle JAVA JDK version 1.7 on the Admin node.



Version 1.7 is required; version 1.7u15 is recommended.

To verify the version of JDK you are running on the Admin node, perform the following steps as both the root and gpadmin users:

```
$ /usr/java/default/bin/java -version
```

The output of this command should contain 1.7 (version number) and JavaHotSpot(TM) (Java version). For example:

```
java version "1.7.0_45"

Java(TM) SE Runtime Environment (build 1.7.0_45-b18)

Java HotSpot(TM) 64-Bit Server VM (build 24.45-b08, mixed mode)
```

If you are not running the correct JDK, download a supported version from the Oracle site at http://www.oracle.com/technetwork/java/javase/downloads/index.html

Install the JDK on the admin node and add it to alternatives as follows:

```
# sudo /usr/sbin/alternatives --install "/usr/bin/java" "java" "/usr/java/jdk1.7.0_xx/bin/java"
3
# sudo /usr/sbin/alternatives --install "/usr/bin/javac" "javac"
"/usr/java/jdk1.7.0_xx/bin/javac" 3
# sudo /usr/sbin/alternatives --config java
```

OpenJDK

Make sure you are not running OpenJDK as your default JDK.

If you are running OpenJDK, we recommend you remove it.

To check for all versions of JDK that are running on your system, as root run:

```
yum list installed | grep jdk
```

An example output from this command is:

```
java-1.6.0-openjdk.x86_64
java-1.7.0-openjdk.x86_64
jdk.x86_64 2000:1.7.0_45-fcs
```

This indicates that there are three versions of JDK installed, two of them are OpenJDK.

To remove all OpenJDK versions, as root, run:

```
yum erase *openjdk*
```

3. **Package Accessibility**. Verify that all packages are available in a local yum repository or that you have yum access to an EPEL yum repository. See Package Accessibility, below.

iptables. Verify that iptables is turned off:

As root:

```
$ chkconfig iptables off
$ service iptables stop
```

4. SELinux. Verify that SELinux is disabled:

As root:

\$ sestatus

If SELinux is disabled, one of the following is returned:

SELinuxstatus: disabled

or

SELinux status: permissive

If SELinux status is *enabled*, you can temporarily disable it or make it permissive (this meets requirements for installation) by running the following command:

As root:

\$ echo 0 > /selinux/enforce



This only temporarily disables SELinux; once the host is rebooted, SELinux will be re-enabled. We therefore recommend permanently disabling SELinux, described below, while running Pivotal HD/HAWQ (however this requires a reboot).

You can permanently disable SE Linux by editing the /etc/selinux/config file as follows:

Change the value for the SELINUX parameter to:

SELINUX=disabled

Then reboot the system.

• Internet Connectivity. You will need an active internet connection to install PCC. This is required to pull the software dependencies for the UI.

Package Accessibility

Pivotal Command Center and Pivotal HD Enterprise expect some prerequisite packages to be pre-installed on each host, depending on the software that gets deployed on a particular host. In order to have a smoother installation it is recommended that each host would have yum access to an EPEL yum repository. If you have access to the Internet, then you can configure your hosts to have access to the external EPEL repositories. However, if your hosts do not have Internet access (or you are deploying onto a large cluster), then having a local yum EPEL repo is highly recommended. This will also give you some control on the package versions you want deployed on your cluster. See Creating a YUM EPEL Respository for instructions on how to setup a local yum repository or point your hosts to an EPEL repository.

For Pivotal Command Center 2.2.x, here is a list of prerequisites that need to either already be installed on the Command Center admin host or on an accessible yum repository:

- httpd
- mod_ssl
- postgresql
- postgresql-devel
- postgresql-server
- compat-readline5
- createrepo
- sigar
- sudo

Run the following command on the admin node to make sure that you are able to install the prerequisite packages during installation.

```
$ yum list <LIST_OF_PACKAGES>
```

For example:

\$ yum list httpd mod_ssl postgresql postgresql-devel postgresql-server compat-readline5 createrepo sigar sudo

If any of them are not available, then you may have not added the repository correctly to your admin host.

For the cluster hosts (where you plan to install the cluster), the prerequisite packages depend on the software you will eventually install there, but you may want to verify that the following two packages are installed or accessible by yum on all hosts:

- nc
- postgresql-devel

For example:

\$ yum list nc postgresql-devel

Installation Instructions

This section provides detailed installation steps. If you are an advanced user you may want to use the more succinct PCC Installation Checklist.

Perform the following installation steps as a root user.



Avoid using hostnames that contain capital letters because Puppet has an issue generating certificates for domains with capital letters.

Also avoid using underscores as they are invalid characters in hostnames.

If you are upgrading from an earlier version of Pivotal Command Center, see Upgrading PCC.

Installing Pivotal Command Center

- 1. Download PCC from the Pivotal Network. PCC is also available from EMC Download Center at https://emc.subscribenet.com.
- 2. As root, create a directory (phd) for your PCC installation on the Admin node:

```
$ mkdir phd
```

3. Copy the Pivotal Command Center tar file to the Admin node. For example:

```
$ scp ./PCC-2.2.x.version.build.os.x86_64.tar.qz host:/root/phd/
```

4. As root, cd to the directory where the Command Center tar files are located and untar them. For example:

```
$ cd /root/phd
  $ tar --no-same-owner -zxvf PCC-2.2.x.version.build.os.x86_64.tar.gz
```

5. Still as root user, run the installation script. This installs the required packages, configures Pivotal Command Center, and starts services.



Important

You must run the installation script from the directory where it was extracted; for example: For example: PCC-2.2.x.version

For example:

```
$ cd PCC-2.2.x.version
$ ./install
```

You will see installation progress information on the screen. Once the installation successfully completes, you will receive an installation success message on your screen.

6. Enable Secure Connections (optional):

Pivotal Command Center uses HTTPS to secure data transmission between the client browser and the server. By default, the installation script generates a self-signed certificate.

Alternatively you can provide your own Certificate and Key by following these steps:

- a. Set the ownership of the certificate file and key file to gpadmin
- b. Change the permission to owner read-only (mode 400)
- c. Edit /etc/httpd/conf.d/pcc-vhost.conf file and change following two directives to point to location of ssl certificate and key, for example:

```
SSLCertificateFile: /usr/local/greenplum-cc/ssl/<servername>.cert
SSLCertificateKeyFile: /usr/local/greenplum-cc/ssl/<servername>.key
```

d. Restart PCC by running:

```
$ service commander restart
```

7. Verify that your PCC instance is running:

```
$ service commander status
```

The PCC installation you just completed includes a CLI (Command Line Interface tool: icm_client). You use this tool to import the PHD services; then you can configure and deploy a cluster using the Command Center User Interface.

You can switch to the gpadmin user (created during installation) for the rest of the installation process:

```
$ su - gpadmin
```

Import Packages

Once you have Pivotal Command Center installed, you can use the <code>import</code> option of the <code>icm_client</code> tool to synchronize the PHD service RPMs and a downloaded JDK package from the specified source location into the Pivotal Command Center (PCC) local yum repository of the Admin Node. This allows the cluster nodes to access the packages during deployment.

If you need to troubleshoot thiS part of the installation process, see the log file located at:

/var/log/gphd/gphdmgr/gphdmgr-import.log

Import JDK

Note that having having JDK 1.7 running on the Admin node is a prerequisite. This step is to import a downloaded JDK package that will be deployed across the cluster.

- Download a supported JDK package from http://www.oracle.com/technetwork/java/javase/downloads/index.html.
- 2. Import the downloaded JDK package to the cluster nodes: As gpadmin, run:

```
$ icm_client import -r <PATH TO JDK>
```

Copy the PHD Service Packages

1. Download the PHD service packages (PHD, and optionally ADS) from the Pivotal Network. The packages are also available from EMC Download Center at https://emc.subscribenet.com.

2. Copy the Pivotal HD, and optionally ADS (HAWQ) tarballs from your initial download location to the gpadmin home directory on the Admin node (home/gpadmin).

```
# If the file is a tar.gz or tgz, use
$ tar zxf PHD-2.0.x-<BUILD>.tar.gz
# If the file is a tar, use
$ tar xf PHD-2.0.x-<BUILD>.tar
# Similarly for the Pivotal ADS tar.gz or tgz file, use
$ tar zxf PADS-1.2.x-<BUILD>.tar.gz
# If the file is a tar, use
$ tar xf PADS-1.2.x-<BUILD>.tar
```

Import PHD Service

1. As gpadmin, import the following tarball for Pivotal HD:

```
$ icm_client import -s <PATH_OF_EXTRACTED_PHD_PACKAGE>
```

For example:

```
$ icm_client import -s PHD-2.0.x-x/
```

Import HAWQ/PXF Services



This is required only if you wish to deploy HAWQ.

1. As gpadmin, import the following tar balls for HAWQ and PXF:

```
$ icm_client import -s <PATH_OF_EXTRACTED_ADS_PACKAGE>
```

For example:

```
$ icm_client import -s PADS-1.2.x-x/
```

You are now ready to configure and deploy a cluster from the Pivotal Command Center UI.

Launching Pivotal Command Center

Launch a browser and navigate to the host on which you installed Command Center. For example:

```
https://CommandCenterHost:5443
```

The Command Center login page is launched in your browser. The default username/password is <code>gpadmin/Gpadmin1</code> (case sensitive).

Next Steps

See Using PCC for details about using the application, including how to change the default password and how to deploy and configure a HD cluster via the Command Center UI.

See the *Pivotal HD Enterprise Installation and Administrator Guide* for instructions for using the command-line interface (CLI) of Pivotal Command Center to deploy and configure a HD cluster.

Starting, Stopping, and Restarting Command Center Services

To stop or restart Command Center services, as root run the following commands on the Pivotal Command Center admin host:

```
$ service commander stop
$ service commander start
$ service commander restart
```



These commands need to be run as root.

Uninstalling Pivotal Command Center

Follow the steps below to uninstall Pivotal Command Center and the Pivotal HD cluster:

- 1. As gpadmin, stop services on all your clusters (See the *Pivotal HD Enterprise Installation and Administrator Guide* for detailed steps).
- 2. As gpadmin, uninstall all your clusters (See the *Pivotal HD Enterprise Installation and Administrator Guide* for detailed steps).
- 3. From the directory where you untarred the Pivotal Command Center, run the uninstall script as root:

```
# cd /root/phd/PCC-2.2.x.version/
```

./uninstall

Upgrading Pivotal Command Center

The following instructions are for upgrading Pivotal Command Center from version 2.1.x to 2.2.



Upgrade Notes

- If you are upgrading to a new version of Pivotal Command Center, make sure you are running compatible versions of Pivotal HD and Pivotal ADS (optional).
- See the latest version of the Pivotal Command Center Release notes for Pivotal Interoperability Matrix.
- We recommend that you always back up any critical data before performing any upgrades.
- As of PCC 2.1.1, there is a new default password, Gpadmin1, and new password rules are being enforced (minimum 8 characters, case-sensitive, 1 uppercase letter, one number). However if you are upgrading to this version of PCC, your original password (by default gpadmin) is maintained.
- JDK 1.7 is a requirement for PCC version 2.2. If you were using JDK 1.6 with an earlier version of PCC, you must install JDK 1.7 before upgrading to PCC 2.2.
- You don't have to stop your cluster to upgrade PCC, it can be either running or stopped.

Follow the steps below to upgrade your Pivotal CC to a newer version:

1. As root, stop the Command Center services:

```
$ service commander stop
```

- 2. Download the new PCC tarball and untar.
- 3. Run the installer from the new PCC location:

```
$ ./install
```

Note that your installation will fail if you do not have the required JDK 1.7 installed.

4. Make sure that nmon is running:

To check nmon status:

\$ service nmon status

To start nmon:

\$ service nmon start

Chapter 4 Using PCC

This section provides an overview of the Pivotal Command Center 2.2 user interface and details about using the application to configure and deploy a Pivotal HD Cluster.

Topics:

- Overview
- Logging In
- Settings
- User Management
- Cluster Status Page
- Configuring and Deploying a Cluster
- · Starting Stopping and Uninstalling the Cluster
- Dashboard
- Cluster Analysis
- MapReduce Job Monitor
- YARN App Monitor
- HAWQ Query Monitor
- Topology
- Logs

Overview

Pivotal Command Center UI is a browser-based application for configuring, deploying, administering, and monitoring Pivotal HD clusters. At a high level, the screens consist of:

Cluster Status Page—Provides status information about any clusters you have configured and deployed.
 Also provides access to the Add Cluster Wizard that allows you to configure and deploy clusters from the UI. See Configuring and Deploying a Cluster for more details.

• Dashboard—Provides an overview of your Pivotal HD cluster. This screen shows at one glance the most important states and metrics that an administrator needs to know about the Pivotal HD cluster.

Cluster Analysis—Provides detailed information about various metrics of your Pivotal HD cluster. This
provides cluster-wide metrics all the way down to host-level metrics.

MapReduce Job Monitor—Provides details about all, or a filtered set of MapReduce jobs.

YARN App Monitor—Provides details about all, or a filtered set of YARN applications.

 HAWQ Queries—When HAWQ (a revolutionary MPP database on Hadoop solution) is deployed on the cluster, Command Center can show the progress of all actively running queries on HAWQ.

 Topology—This screen shows you what roles have been installed on each host. You can also add and remove slaves to/from the cluster from this screen.

 Logs—This screen displays system logs based on filter criteria you select such as log levels, host name, time period, and so on.

Status indicators

Throughout the user interface the following indicators are used to indicate the status of nodes:

· Green: Succeeded

Blue: Running

Grey: Stopped/Pending

Red: Killed/Failed

Logging In

Launch a browser and navigate to the host on which you installed Command Center. For example:

https://CommandCenterHost:5443

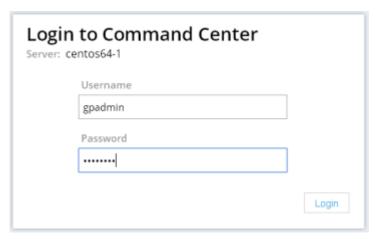
The Command Center login page is launched in your browser. The default username/password is <code>gpadmin/Gpadmin1</code> (case-sensitive).

To change the default port (5443), update the port settings in the following file:

/usr/local/greenplum-cc/config/app.yml

Login Screen

The first time you launch the Command Center UI, a login screen appears showing the hostname of this instance of Pivotal Command Center.



The default admin username/password is <code>gpadmin/Gpadmin1</code> (case-sensitive). You can change this password via the Settings menu.

Passwords are case-sensitive and must be at least 8 letters long and contain 1 upper-case letter and 1 number.

Once you have entered a valid username/password, click the Login button to launch the Command Center UI.

Settings

Once you have logged in, you can click the gear icon in the upper right corner of the screen from any PCC page to display the **Settings** menu.



From the settings menu you can select one of:

- About. Select this to display version information about this instance of PCC
- Cluster Status. Select this option to go back to the Cluster Status page to view the list of available clusters.
- Users. Select this option to add/edit user information. There are two types of users, administrative and non-administrative (read-only); you will only see this Users option if you are an administrative user. See Users below.
- Change Password. Click this to change your password. See Passwords below.
- Edit Profile. Click this to change your profile. See Users below
- Logout. Select this option to logout from this instance of PCC.

User Management

There are two types of users, administrative (super user) and non-administrative (read-only).

If you are an administrative user, you can add users, edit your own or other users' profiles, change your own or other users' passwords, and delete users.



Non-administrative users are read-only users; their actions are limited to read-only actions.

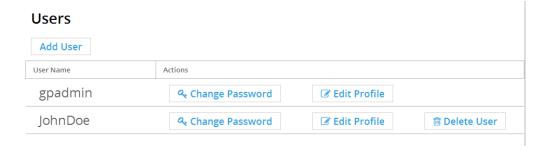


For many user management actions, you will be asked to Re-authenticate to Continue. In all cases, this means you need to enter the password of the current user performing the action to continue.

Users

You will only see the **Users** option in the **Settings** menu if you are an administrative user.

Select **Users** from the **Settings** menu to see a screen listing all current users:



Adding a User

1. Click Add User to create a new user.

A New User form appears. Enter the following information about the new user:

- First Name.
- Last Name.
- User Name. This must be a unique name.
- Email. Must be a valid email address.
- Administrator. Check this box if the new user is to have administrative privileges.
- Password. Enter a password for the new user. This must meet the minimum password requirements.
 Passwords are case-sensitive and must be at least 8 letters long and contain 1 upper-case letter and 1 number.
- Confirm Password. Confirm the new user's password.
- 2. Click OK.

A dialog box appears asking you to **Re-authenticate to Continue**.

- 3. Enter the password of the current user (the administrative user who is creating the new user).
- 4. Click OK.

You are returned to the Users screen where you should now see your new user(s) in the list.

Deleting a User

Only administrative users can delete users and the default gpadmin user can never be deleted. To delete a user:

- 1. Select Users from the Settings menu
- Click the **Delete User** link adjacent to the user you want to delete.A dialog box appears asking you to Re-authenticate to Continue.
- 3. Enter the password of the current user (the administrative user who is deleting the user).
- 4. Click OK.

You are returned to the Users screen where the user you just deleted should no longer appear in the list.

Profiles

Your profile includes: First Name, Last Name, User Name, Email Address, and User Type (administrative or not); it does not include your password.

All users can edit their own profiles by:

• Selecting Edit Profile from the Settings menu.

Administrative users can edit their own passwords and also the profiles of other users by:

Selecting Users from the Settings, then clicking the Edit Profile link adjacent to the user whose profile
they wish to edit.

After making edits to any profile, a dialog box appears asking you to **Re-authenticate to Continue**. Enter the password of the current user making the edits to continue, then click **OK**.

Passwords

Passwords are case-sensitive and must be at least 8 letters long and contain 1 upper-case letter and 1 number.

All users can change their own passwords by:

• Selecting Change Password from the Settings menu.

Administrative users can change their own passwords and also the passwords of other users by:

 Selecting Users from the Settings menu, then clicking the Change Password link adjacent to the user whose password you wish to change.

After changing any password, a dialog box appears asking you to **Re-authenticate to Continue**. Enter the password of the current user making the edits to continue, then click **OK**.

Cluster Status Page

Once you have launched Command Center, the initial screen you see is the Cluster Status screen. This displays a list of available clusters to monitor, the status of each cluster (**started**, **stopped**), and a list of services running on that cluster (Hive, Mahout, and so on).



The Add Cluster and Actions buttons are only visible to administrative users.



From this page you can:

- Administrative users only: Click Add Cluster to launch the Add Cluster Wizard.
- Click the cluster name in the table to view the Dashboard for that cluster.
- From any point within Command Center UI, you can always select a different cluster by using the **Select**Cluster drop-down menu in the upper right corner of the screen.
- Administrative users only: You can either **Start**, **Stop**, or **Uninstall** a cluster. Depending on the state of the cluster, some of these buttons will be enabled while others are disabled.

Configuring and Deploying a Cluster



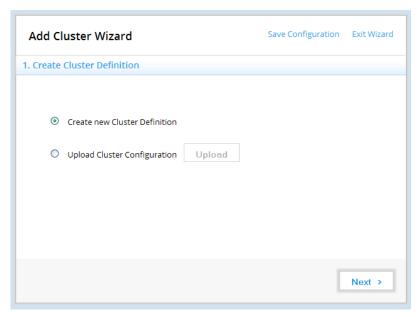
Before you can configure and deploy a cluster, make sure you have already installed and enabled the PHD Services (see Install PHD Services and Enable the PHD Services).

After you have logged in to Pivotal Command Center, the Cluster Status page appears. From here, if you are an administrative user, you are able to launch the Add Cluster Wizard that enables you to configure and deploy a Pivotal HD Cluster.

As you move through the wizard, the right hand pane displays where you are in the deployment process:

Adding a Cluster

1. Click **Add Cluster**. The Add Cluster Wizard opens:

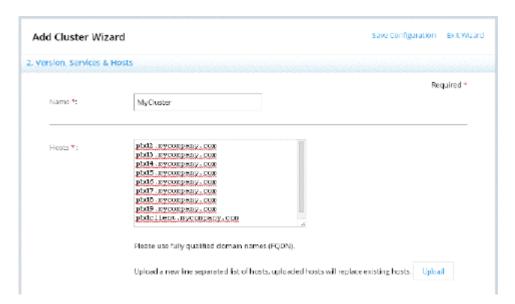


The Wizard allows you to create a new configuration from scratch or upload and edit any existing configuration. The Summary panel along the right shows you the progress of your configuration and deployment.

2. Create Cluster Definition. Either:

- a. If you are configuring a new cluster, select Create a new Cluster Definition then click Next.
- b. If you want to edit an existing cluster; select Upload Cluster Configuration, click Upload, then navigate to the clusterConfig.xml file that you wish to edit; then click Next. In this case, the following fields in the Wizard will be populated with the cluster definition properties of that clusterConfig.xml file you just uploaded. Follow the instructions below to edit those values.

3. Versions, Services and Hosts:



\triangle

Notes

- Hosts can be entered individually, newline-separated; or can be expressed in a range, for example host[1-5].yourdomain.com. They can also be expressed in multiple ranges, for example host [1-3].subdomain[1-2].yourdomain.com. Any hosts expressed in ranges are expanded during host verification. Hosts that do not exist within a specified range will be ignored, so you can specify a wide range and only those hosts that are available within that range will be added.
- If you are editing an existing configuration, some if not all of these fields will be pre-populated. Edit where appropriate.
- You need to scroll down to view all the fields on this screen. The **Next** button will not be active until you have entered all the required fields.

Enter the following information:

- Name: Required. Enter a name for this cluster. Special characters are not supported.
- Hosts: Required. Enter a new line-separated list of FQDN host names. You can also click Upload to
 use a text file containing a new line-separated list of host names.
- Root Password: Required. Enter the root password.
- GP Admin Password: Required. Enter the gpadmin user password. PCC creates this user on all nodes.
- JDK Path: Enter the JDK filename (not the absolute path). For example: jdk-7u51-linux-x64.rpm.

Note: JDK 1.7.0_15 (min) is a prerequisite.

If not already installed, you should download it then install using icm_client import -r (see Installing PCC for more details)

- Setup NTP: Check this box if you want to set up NTP (Network Time Protocol).
- Disable SELinux: Check this box if you want to disable SELinux. Recommended.
- Disable IPTables: Check this box if you want to disable IPTables. Recommended.
- Run ScanHosts: Leave this box checked if you want to run scanhosts. The scanhosts command verifies the prerequisites for the cluster node have been met and provides a detailed report of any missing prerequisites. Running this command ensures that clusters are deployed smoothly and is strongly recommended.

Click Next.

4. Host Verification:

The Host Verification page opens. This step may take a few minutes, it verifies connections to the hosts you just set up. Once the Eligibilty field changes from Pending, to Eligible for all hosts, you can click **Next.** You will see any error and informational messages displayed in the comments fields.



If you specified hosts using ranges, they will be expanded at this point.

5. Topology:

This is the section where you specify the roles to be installed on the hosts. For example, you can specify where your hadoop namenode, data node and so on, should be installed. Note that all mandatory roles should have at least one host allocated.

Each service has its own section on this page; you can use the top menu options as shortcuts to those sections on the page, or simply scroll down to each section.

Notes

 You need to click Enter or Tab before each field is accepted. Once you enter the text and click Enter or Tab, the text will change appearance and appear enclosed in a box, as shown in the figure below. The entry on the left has been accepted, the entry on the right has not.

```
centos62-2 * invalidhost.com
```

- Hosts can be specified in ranges, see the notes for Versions, Services and Hosts, for more information.
- At any point during this stage you can click Save Configuration at the top right of the page.
 This saves the configuration file and downloads it. Once saved, a link to the configuration file appears at the bottom of the page. Click that link to open and view the clusterConfig.xml file. You cannot edit this xml file directly.

These are the roles that need to have installation nodes defined:

- CLIENT: ICM installs Pig, Hive, HBase, and Mahout libraries on this host.
- HDFS: Name Node, Secondary Name Node, Data Nodes
- YARN: Resource Manager, History Server, Node Managers
- Zookeeper: Zookeeper Server
- HBase: Hbase Master, HBase Region Servers.
- Hive: Hive Master, Hive Metastore
- HAWQ: Primary Node, Secondary Node, HAWQ Segment Nodes
- USS: Name Node and Catalog
- PXF: No hosts to configure. Installed on the client host.
- Mahout: No hosts to configure. Installed on the client host.
- Pig: No hosts to configure. Installed on the client host.

Click **Next** once you have finished role-mapping.

6. Cluster Configuration:

This page displays a list of all configuration files that define this cluster; the clusterConfig.xml (to edit service configuration global values) as well as the service specific configuration files.

All these configuration files are already populated with the values you have already entered; or with default values.

Click any file name to open that configuration file in an editor and enter/edit values.

If you make any changes, click **Save** to return to the Cluster Configuration page.

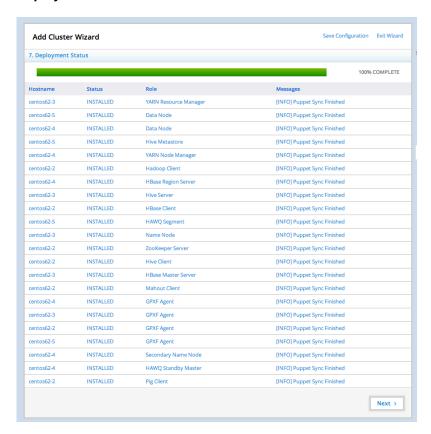
Once you have completed all your edits, click Next.

7. Validation:

If the configuration has errors they will be displayed here; otherwise you will see post-deployment instructions.

Click **Deploy**

8. Deployment Status:



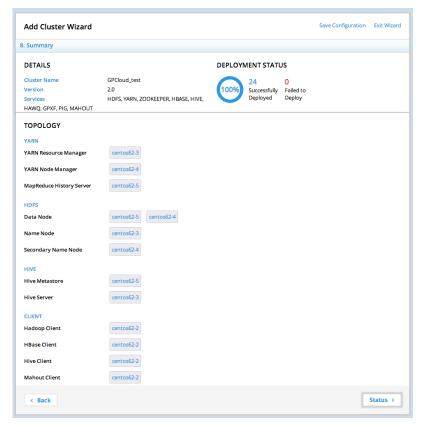
This screen shows the progression of the deployment. Information displayed includes:

- Hostname
- Status
- Role
- Messsages

Once the deployment is complete, click Next.

9. Summary

Once your cluster has successfully deployed, you can view a summary of the cluster, as shown here:



10. Return to the Cluster Status page:

Once you have reviewed this summary, click **Status**, to return to the Cluster Status page. Your new cluster will be listed on this page, with the status of **installed**.

If you are deploying HAWQ you need to initialize and configure HAWQ after you start the cluster. See Initializing and Configuring HAWQ.

Next steps: Starting the Cluster then Initializing and Configuring HAWQ.

Starting the Cluster

To start your cluster; click **Actions**: **Start** on the Cluster Status page.

Initializing and Configuring HAWQ

For HAWQ users:

ssh to the HAWQ master, then as gpadmin, run the following:

su - gpadmin

```
# source /usr/local/hawq/greenplum_path.sh
```

```
# /etc/init.d/hawq init
```

Add the IP address of your instance of Command Center to HAWQ's pg_hba.conf file, for example:

```
vi /data1/master/gpseg-1/pg_hba.conf
```

host all gpadmin <Command_center host IP>/24 trust

Then restart HAWQ:

```
# /etc/init.d/hawq restart
```

You have now completed your cluster configuration and deployment.

See the following sections of this document for details about using the PCC UI to administer and monitor your cluster.

Starting Stopping and Uninstalling the Cluster

These functions are only available to administrative users, also, depending on the state of the cluster, some of these buttons will be enabled while others are disabled.

Starting

To start a cluster:

From the Cluster Status page locate the cluster you want to start, then click the Actions: Start button.

Stopping

To stop a running cluster:

From the Cluster Status page locate the cluster you want to stop, then click the **Actions**: **Stop** button.

Uninstalling

Only stopped clusters can be uninstalled.

To uninstall a cluster:

- 1. From the Cluster Status page locate the cluster you want to uninstall, then click the **Actions**: **Uninstall** button.
- 2. An Uninstall confirmation dialog appears. Click **Yes** to confirm you want to uninstall this cluster; **No** to cancel the operation and return to the Cluster Status page.
- 3. Once you confirm you want to continue with the uninstall operation, a Preserve Cluster State dialog appears. You need to specify whether you want to preserve the data for this cluster. Your options are: **Cancel**. This cancels the operation and returns you to the Cluster Status page.
 - **Yes**. This preserves the data for this cluster before uninstalling the cluster and returning you to the Cluster Status page.
 - No. This uninstalls the cluster without preserving the data; then returns you to the Cluster Status page.

Dashboard

The dashboard gives you a high level view of a cluster at a glance. You are able to view the status of the most important cluster services, such as HDFS and YARN, and allows you to start and stop each service individually. It also shows you how the most important cluster metrics are trending in a visual way.

The graphs provide a unified view of the state of your system. They are also useful in detecting outliers and pinpointing specific problems that may be present in your system.



The right side of the Dashboard displays the state of the following services, provided they have been deployed for this cluster:

HDFS

For HDFS, the dashboard provides the following information/functionality:

- The status of HDFS. You can use the Actions dropdown menu to Start/Stop HDFS depending on its status.
 - When the last NameNode checkpoint occurred.
 - The percentage of cluster storage being used by HDFS and how much is free.
 - The number of DataNodes that are up and whether they are running normally or with problems.
 - The Actions dropdown menu allows you to Rebalance (redistribute your data across the cluster) your cluster, and to View Rebalancer Log.

The only Action option available to non-administrative users is to View Rebalancer Log.

If High Availability (HA) is enabled for your cluster, you will see the status of two NameNodes here.

YARN

For YARN, the dashboard provides the following information:

- The status of YARN. You can use the Actions dropdown menu to Start/Stop YARN depending on its status (not available for non-administrative users).
 - Whether or not the YARN History Server is running.

Note: The History Server stores a history of the mapreduce jobs run on the cluster.

The number of NodeManagers that are running.

HBase

For HBase, the dashboard provides the following information:

 The status of the HBase master. You can use the Actions dropdown menu to Start/Stop HBase depending on its status (not available for non-administrative users).

Zookeeper

For Zookeeper, the dashboard provides the following information:

- The status of Zookeeper. You can use the Actions dropdown menu to Start/Stop Zookeeper depending on its status (not available for non-administrative users).
 - The Status of the Zookeeper server host.

Hive

For Hive, the dashboard provides the following information:

 The status of Hive. You can use the Actions dropdown menu to Start/Stop Hive depending on its status (not available for non-administrative users).

The Dashboard also provides metrics about:

- Mapreduce Slot Utilization
- Namenode RPC Times
- Hadoop Datanodes Average CPU

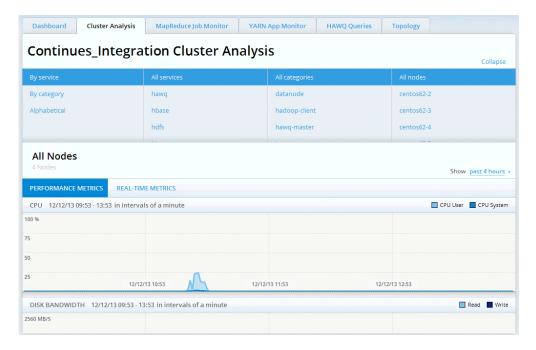
- Hadoop Datanodes Average Bandwidth
- Namenode Operations Per Second
- Hadoop Datanodes Average Disk Bandwidth
- Hadoop Datanodes Average Memory
- Mapreduce Jobs By Status

Cluster Analysis

The Cluster Analysis screen provides detailed metrics on your Pivotal HD cluster.

It provides cluster-wide metrics all the way down to host-level metrics. It provides Hadoop-specific metrics, as well as system metrics that you can drill down to if needed.

The Cluster Analysis screen displays the same data that is shown in the dashboard but in greater detail.



By default the Cluster Analysis screen displays the metrics for all services, all categories, and all nodes. You can filter the information displayed by combinations of the following filters:

By Service. Metrics can be filtered by services such as HDFS, YARN, or HAWQ.

By Category. Metrics can be filtered by categories such as:

- namenode
 - secondarynamenode
 - datanode
 - yarn-resourcemanager
 - yarn-nodemanager
 - mapreduce-historyserver
 - hawq-master
 - hawq-segment

Alphabetically. Metrics can be filtered alphabetically.

Based on the filters you select, the lower part of the Cluster Analysis screen provides detailed graphs that display data related to:

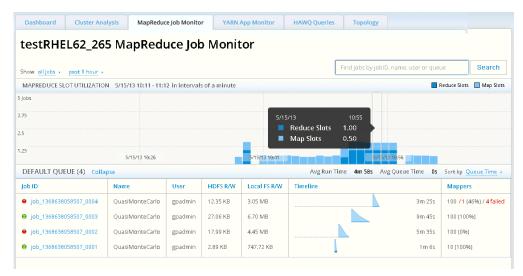
- CPU
- Disk Bandwidth
- Network Bandwidth
- Memory
- Load
- Swap Usage
- Swap I/O
- Network Operations
- Disk Operations

You can view either the **Performance Metrics**, which show the cluster/node utilization over-time, the **Real-time Metrics** which show the current metrics in real-time, or **Storage Metrics**, which show metrics about cluster storage.

If you select Cluster Analysis for **All Nodes** (the default), the Trending Metrics graph for the cluster is displayed.

MapReduce Job Monitor

The Job Monitor screen tracks the MapReduce jobs that are executed in the Pivotal HD cluster when the YARN MapReduce service is running. It provides details about all, or a filtered set of MapReduce jobs.



The MapReduce jobs displayed can be filtered by state and/or time range.

By state:

- all jobs (set by default)
 - · currently pending jobs
 - · currently running jobs
 - · succeeded jobs
 - · failed jobs
 - killed jobs
 - · error state jobs

By time range: By selecting a preset time range in hours, weeks, months, year, or by specifying a custom time range.

The MapReduce jobs can also be filtered by searching for values for the following:

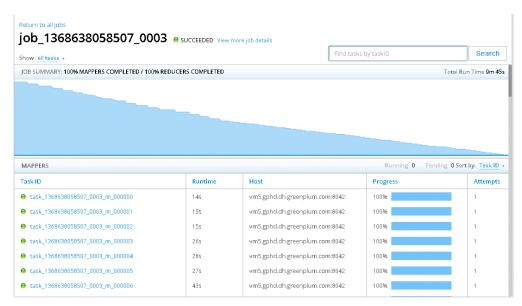
- jobID
- name
- user
- queue

Enter your search value in the search bar in the following format: searchKey=searchValue, where searchKey is one of **jobID**, **name**, **user**, or **queue**.

These are substring searches. For example: **jobID=1363920466130** will locate a job with **jobID=job_1363920466130_0002**

Job Details

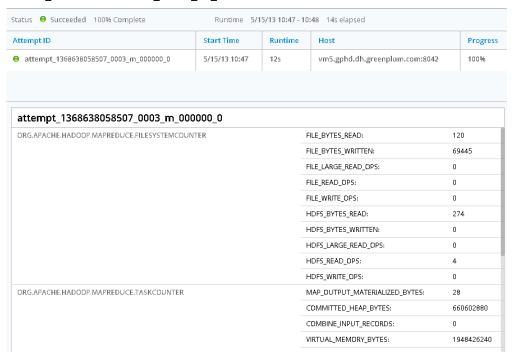
When you click on any of the jobs in the Job Monitor more details of the job are shown.



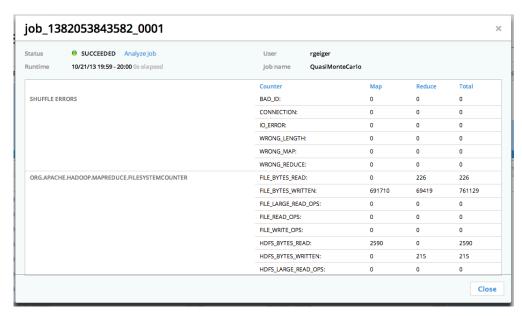
This screen displays all the tasks that are have been allocated for the selected job and their progress. You can see the mapper and the reducer tasks separately. In the above screen capture, the bars in the JOB SUMMARY section represent the two Mapper tasks that have run, one took 19 seconds, the other, 20 seconds.

Clicking on each task ID will show even more details about that particular task. You can also filter on a particular task ID in the search bar.

task_1368638058507_0003_m_000000



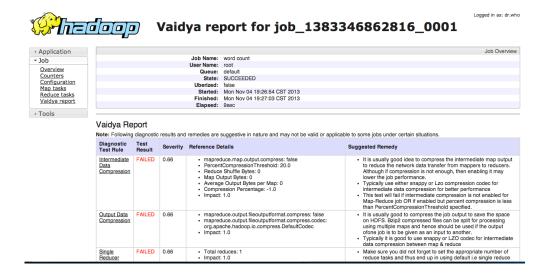
To see job related counters click on View more job details next to the job ID:



Click the **Analyze Job** link adjacent to the Status field to open a Vaidya report about the selected job, as shown below:



This capability is beta and will be improved in coming releases.



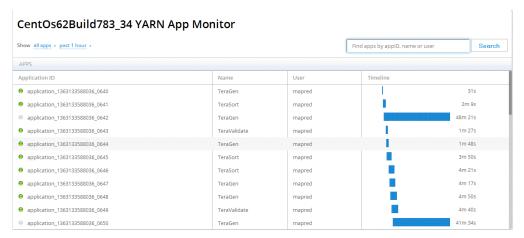
About Vaidya

Vaidya is a diagnostic tool installed with PHD for Map/Reduce jobs. After a job is executed successfully, it uses a job history log and job configuration information to identify any performance or scalability problems with the job. Upon execution, it provides a job analysis report indicating specific problems with the job along with the remedy to correct them.

For more information about Vaidya, see the PHD Enterprise Stack and Tool Reference Guide.

YARN App Monitor

The YARN App Monitor screen tracks YARN applications that are executed in the Pivotal HD Cluster.



The YARN applications displayed can be filtered by category and/or time range:

- By Category:
 - all apps (set by default)
 - · currently pending apps
 - currently running apps
 - succeeded apps
 - · failed apps
 - killed apps
 - error state jobs
- By Time Range: By selecting a preset time range in hours, weeks, months, year, or by specifying a custom time range.

The YARN applications can also be filtered by the following fields by entering it in the search bar in the following format: searchKey=searchValue:

- appID
- name
- user

These are substring searches. For example: appID=1363920466130 will locate the application with appID=application_1363920466130_0002

HAWQ Query Monitor

The HAWQ Query monitor is only displayed when HAWQ is installed on the cluster.

This screen displays all active queries running on the HAWQ cluster:



In this release, this screen only displays active queries as can be seen when you run:

SELECT * FROM pg_stat_activity;

on the HAWQ cluster.

Click on a Query ID to get the syntax of that query:



Topology

This screen shows you what roles have been installed on each host. You can also add and remove slaves to/from the cluster via this screen.



A node is considered a slave if it contains the following services: datanode, hbase-regionserver , yarn-nodemanager.



Topology Actions



This functionality is restricted to administrative users.

Adding Slaves to the Cluster

Click on the Add Slaves to the Cluster option from the Topology Actions menu. An Add Slaves dialog appears.

This dialog lists all current nodes. Enter the slave nodes you want to add either individually, or as ranges, for example Node[1-9], provide the Root Password, then click Add Slaves to Cluster.

You are warned that new slaves added to the cluster are not started automatically. To start a slave node, return to the Dashboard and navigate to the role defined for that node, then select Action > Start <service name>, for example Action > Start Hive.



Newly added slaves are not started automatically.

- If you type a node name twice, that name will flash yellow until one is deleted.
- If you type the name of a node that already exists, that name becomes highlighted and you are shown an error message warning you that:

Highlighted hosts are duplicates and will not be added.

 If you close the dialog before the operation is finished, it continues in the background. Refresh the Topology page to see if it has successfully completed.

Removing Slaves from a Cluster

Click on the **Remove Slaves from the Cluster** option from the **Topology Actions** menu. A **Remove Slaves** dialog appears; this dialog lists all current slave nodes

You are warned that prior to removal, services on slave nodes must be stopped. To stop services on a slave node, return to the Dashboard and navigate to the role defined for that node, then select **Action > Stop** <service name>, for example **Action > Stop Hive**.

Enter the slave nodes you want to remove either individually, or as ranges, for example Node[1-9], then click Remove Slaves from Cluster.



- Prior to removal, services on slave nodes must be stopped .
- You do not need to provide a password to remove slave nodes.
- This text field auto-completes with nodes from the cluster.
- If you try and enter a node that does not exist, the text field does not become active and you are not able to perform the remove operation.
- If you close the dialog before the operation is finished, it continues in the background. Refresh the Topology page to see if it has successfully completed.

Logs

This screen displays all system logs.

Filtering Logs

You can use the dropdown menus to filter what logs are displayed based on:

- Logs Levels: You can select to display all logs (the default), or filter by one of the following log levels: debug, errors, fatal, info, trace, warnings
- Hosts: You can select to display logs for all hosts (the default), or can filter by hostname
- Roles: You can select to display logs for all roles (the default), or can filter by role name
- Time: You can select the time period over which you want to view logs; either for the Past 1 hour, or for one of the other time periods available (up to Past 8 weeks)

Search by Keyword

You can also further filter the logs displayed by searching the logs by keyword.

Enter any single keyword (any single string, including numbers) in the search box, then click Search.

Note that your keyword search is only applied to the set of logs displayed based on any filters you chose from the dropdown menus, described above.

Viewing Logs

The results of your filters, if any, are displayed on the screen. Often there will be too many to display on one page, use the Prev/Next/Page numbers at the top of the screen to navigate through them.

Each entry displays the time stamp, a log message, a time stamp, the host and the role, for example:

2014-02-25 23:05:51,137 WARN org.apache.hadoop.yarn.server.nodemanager.containermanager.AuxServices: The Auxilurary Service named 'mapreduce_shuffle' in the configuration is for class class org.apache.hadoop.mapred.ShuffleHandler which has a name of 'httpshuffle'. Because these are not the same tool...

yarn-nodemanager centos62-4 Show Logs | Show Details

2014-02-25 23:05:47,585 - WARN [NIOServerCxn.Factory:0.0.0.0/0.0.0.0:2181:NIOServerCnxn@347] - caught end of stream exception EndOfStreamException: Unable to read additional data from client sessionid 0x1446b2ad66b0001, likely client has closed socket at org.apache.zookeeper.server.NIOServerCnxn.d...

zookeeper-server centos62-2 Show Logs | Show Details

If the log message is too long to display on the screen it is truncated and a **Show Details** link appears below the message; click to display the entire message.

Click **Show Logs** to display all contents of that log file based on the time stamp. You can see more log messages from before and after that time stamp by clicking the **Fetch More Logs Above/Fetch More Logs Below** links at the top and bottom of the page, respectively.

Chapter 5 Creating a YUM EPEL Repository

Pivotal Command Center and Pivotal HD Enterprise expect some prerequisite packages to be pre-installed on each host, depending on the software that gets deployed on a particular host. In order to have a smoother installation it is recommended that each host would have yum access to an EPEL yum repository. If you have access to the Internet, then you can configure your hosts to have access to the external EPEL repositories. However, if your hosts do not have Internet access (or you are deploying onto a large cluster), then having a local yum EPEL repo would be highly recommended. This will also give you some control on the package versions you want deployed on your cluster.

Following are the steps to create a local yum repo:

- 1. Mount the RHEL/CentOS DVD on a machine that will act as the local yum repo.
- 2. Install a webserver on that machine (for example. httpd), making sure that HTTP traffic can reach this machine.
- 3. Install the following packages on the machine:

```
yum-utilscreaterep
```

4. Go to the directory where the DVD is mounted and run the following command:

```
createrepo .
```

5. Create a repo file on each host with a descriptive filename in the

```
/etc/yum.repos.d/ directory of each host (for example, CentOS-6.1.repo) with the following
contents:[CentOS-6.1]
name=CentOS 6.1 local repo for OS RPMS
baseurl=http://172.254.51.221/centos/$releasever/os/
$basearch/
enabled=1
gpgcheck=1
gpgkey=http://172.254.51.221/centos/$releasever/os/$basearch/RPM-GPG-KEY-CentOS-
```

6. Validate that you can access the local yum repos by running the following command:

```
yum list
```

PCC User Guide **Pivotal Command Center**

Chapter 6 Command Line Reference

This section provides descriptions and syntax for the command line operations you can perform.

Topics:

- Backup and Restore
 - Backup
 - Restore



You need to perform command line operations on the Admin node.

Backup and Restore

You can backup data on the admin node where PCC is installed. Having the backup allows you to restore the admin node and PCC to a given state in case of failures or data corruption.



Backup and restore operations should be performed as root.

Backup

Run the backup script to copy all configuration files and all data in the PCC database into a backup file on a local disk. Once the tar file is created you should copy it off the admin node to a different storage to prevent it from being lost if the admin node fails.



You should backup your data each time you make a configuration or topology change for your cluster or if you add or delete clusters. Trying to restore old backup on the admin node when the cluster topology or configuration has been changed since the backup will result in inconsistent configuration and potentially unusable cluster.

To perform the backup run the following script:

```
# /usr/lib/gphd/gphdmgr/bin/gphdmgr_backup.sh
```

The script does not require any input and produces the backup file phdmgr_backup_[timestamp].tar.gz in the current working directory, where timestamp is the number corresponding to the current system time.

Restore

Run the restore script to replace all configuration files and all data in the PCC database from a backup file on a local disk. There are two restore scenarios:

- Restore data on a new node. If you had a backup from the old admin node and you need to provision a new admin node, install PCC on the new admin node, then restore data from the backup produced on the old admin node. Hostnames, IP addresses and system configuration of the old and new nodes have to be identical.
- Restore data on the same node. You can restore data from the backup file on the same admin node where the backup was produced. Make sure the PCC version has not been updated since the backup was made.



A The restore procedure replaces all your configuration files and database data with the data from the backup file, so use with caution.

To perform the restore run the following script:

```
# /usr/lib/gphd/gphdmgr/bin/gphdmgr_restore.sh phdmgr_backup_[timestamp].tar.gz
Your data will be permanently removed. Do you want to continue? y
```

The script takes the name of the backup file as an argument and asks the user for the confirmation before proceeding with the restore. The script takes care of stopping the PCC services, performing the restore, and restarting the services.